

Assignment 1

Due Date: October 4th, 2023, 10:59:00 pm

Total: 160 marks

General Instructions:

- You are allowed to work directly with one other person to discuss the questions. However, you are still expected to write the solutions/code/report in your own words; i.e. no copying. If you choose to work with someone else, you must indicate this in your assignment submission. For example, on the first line of your report file (after your own name and information, and before starting your answer to Q1), you should have a sentence that reads: *At the end of this assignment, I worked together with my classmate [name] and we have written the solutions/code/report.*
- Your submission should be a zip file (containing the solutions/code/report, the MATLAB script, and the answers to the specific questions), with the answers to the questions and discussion to be submitted to MarkUs directly.
- Submit documents and code files separately. Please store all files in a folder named 'Assignment 1' and then submit the files to MarkUs. Please create a **README.txt** file (inside the folder) that describes the files.
- Do not worry if you are not able to submit multiple times. You can submit multiple times.



Part I: Theoretical Questions

[Question 1] Convolution (10 marks)

[1.a] (5 marks) Calculate and plot the correlation **and** the convolution of $x[n]$ and $h[n]$ specified below:

$$x[n] = \begin{cases} 2 & -2 \leq n \leq 4 \\ 0 & \text{otherwise} \end{cases} \quad h[n] = \begin{cases} 1 & -3 \leq n \leq 1 \\ 0 & \text{otherwise} \end{cases} \quad (1)$$

[1.b] (5 marks) Calculate and plot the correlation **and** the convolution of $x[n]$ and $h[n]$ specified below:

$$x[n] = \begin{cases} 2 & -2 \leq n \leq 4 \\ 0 & \text{otherwise} \end{cases} \quad h[n] = \begin{cases} 2 - |n| & -3 \leq n \leq 1 \\ 0 & \text{otherwise} \end{cases} \quad (2)$$

[Question 2] Polynomial Multiplication and Convolution (5 marks)

Vectors can be used to represent polynomials. For example, 3^{rd} -degree polynomial $(a_3x^3 + a_2x^2 + a_1x + a_0)$ can be represented by vector $[a_3, a_2, a_1, a_0]$.

If \mathbf{u} and \mathbf{v} are vectors of polynomial coefficients, prove that convolving them is equivalent to multiplying the two polynomials they each represent.

Hint: You need to assume proper zero-padding to support the full-size convolution.

[Question 3] Laplacian Operator (10 marks)

The Laplace operator in d -dimensional Euclidean space, defined as the divergence of the gradient of a twice-differentiable real-valued function,

where the latter notation

Now, consider a 2D image grid I with partial derivatives associated with the image grid. In other words, show that ΔI

Hint: Start by using the chain rule.



dimensional Euclidean space, defined as the divergence of the gradient of a twice-differentiable real-valued function,

I_{yy} . Here the second partial derivatives are associated with the image grid. In other words, show that ΔI

(x, y) . Then use the chain rule.

[Question 4] Computational Complexity (5 marks)

Assume that we have a convolution implementation $J = \text{conv2}(F, I)$ that takes two images ($F_{k \times k}$ and $I_{n \times n}$) and returns the output in $O(k^2 n^2)$ time. Given an image $I_{n \times n}$ and two filters $F_{k \times k}$ and $G_{k \times k}$, we want to compute $G * (F * I)$. Is it more efficient to call $\text{conv2}(G, \text{conv2}(F, I))$ or $\text{conv2}(\text{conv2}(G, F), I)$? Briefly justify your answer.

[Question 5] Image Pyramids (10 marks)

In Gaussian pyramids, the image at each level I_k is constructed by blurring the image at the previous level I_{k-1} and downsampling it by a factor of 2. A Laplacian pyramid, on the

other hand, consists of the difference between the image at each level (I_k) and the upsampled version of the image in the next level of the Gaussian pyramid (I_{k+1}).

Given an image of size $2^n \times 2^n$ denoted by I_0 , and its Laplacian pyramid representation denoted by L_0, \dots, L_{n-1} , show how we can reconstruct the original image, using the minimum information from the Gaussian pyramid. Specify the minimum information required from the Gaussian pyramid and a closed-form expression for reconstructing I_0 .

Hint: The reconstruction follows a recursive process; What is the base case that contains the minimum information?

Hint: Express the output of a network as a function of its inputs and its weights of layers.

[Question 6] Backpropagation (10 marks)

Consider a neural network

where x_i denotes input

logistic function:

Suppose the loss function is the squared error loss, i.e. $L(y, \hat{y}) = (y - \hat{y})^2$. Assume the

loss, i.e. $L(y, \hat{y}) =$

(w_1, w_2, w_3)

$(0, 0.3)$

[3.a] (5 marks) Draw the computational graph for the network. Define appropriate intermediate variables and their derivatives.

Define appropriate intermediate variables and their derivatives. Break the function into smaller components.)

[3.b] (5 marks) Given an input data point $(x_1, x_2, x_3, x_4) = (-1.1, 1.9, -1.5, 2.0)$ with true label of 0.0, compute the partial derivative $\frac{\partial L}{\partial w_4}$, by using the back-propagation algorithm. Indicate the partial derivatives of your intermediate variables on the computational graph. Round all your calculations to 4 decimal places.

Hint: For any vector (or scalar) \mathbf{x} , we have $\frac{\partial}{\partial \mathbf{x}}(\|\mathbf{x}\|_2^2) = 2\mathbf{x}$. Also, you do not need to write any code for this question! You can do it by hand.

[Question 7] CNN FLOPs (10 marks)

In this problem, our goal is to estimate the computation overhead of CNNs by counting the FLOPs (floating point operations). Consider a convolutional layer C followed by a max pooling layer P . The input of layer C has 50 channels, each of which is of size 12×12 . Layer C has 20 filters, each of which is of size 4×4 . The convolution padding is 1 and the stride is

2. Layer P performs max pooling over each of the C 's output feature maps, with 3×3 local receptive fields, and stride 1.

Given scalar inputs x_1, x_2, \dots, x_n , we assume:

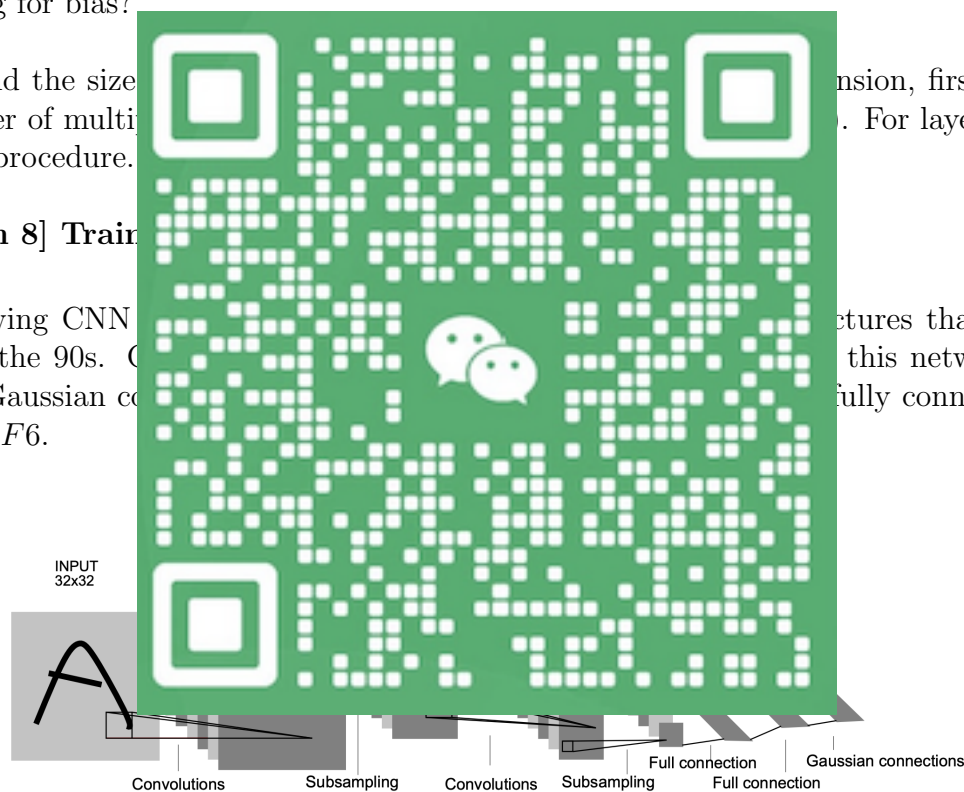
- A scalar multiplication $x_i \cdot x_j$ accounts for one FLOP.
- A scalar addition $x_i + x_j$ accounts for one FLOP.
- A max operation $\max(x_1, x_2, \dots, x_n)$ accounts for $n - 1$ FLOPs.
- All other operations do not account for FLOPs.

How many FLOPs layer C and P conduct in total during one forward pass, with and without accounting for bias?

Hint: Find the size of the input image, the number of multipliers, and the number of additions. For layer P , follow the same procedure.

[Question 8] Training

The following CNN architecture was presented in the 90s. Consider the input image that was presented in this network. Note that the Gaussian connections are similar to $F6$.



[Question 9] Logistic Activation Function (10 marks)

For backpropagation in a node with logistic activation function, show that, in order to compute the gradient, as long as we have the output of the node, there is no need for the input.

Hint: Find the derivative of a neuron's output with respect to its inputs.

Part II: Implementation Tasks (80 marks)

In this question, we train (or fine-tune) a few different neural network models to classify dog breeds. We also investigate their dataset bias and cross-dataset performances. All the tasks should be implemented using Python with a deep learning package of your choice, e.g. PyTorch or TensorFlow.

We use two datasets in this assignment.

1. Stanford Dogs Dataset
2. Dog Breed Images

The Stanford Dogs Dataset (SDD) contains over 20,000 images of 120 different dog breeds. The annotations available for this dataset include class labels (i.e. dog breed name) and bounding boxes. In this assignment, we will only use a small portion of the dataset to train your models on Colab. Dog Breed Images (DBI) contains images of 10 different dog breeds.

To prepare the data:

1- Download both datasets and identify the breeds that appear in both datasets:

- Bernese mountain dog
- Border collie
- Chihuahua
- Golden retriever
- Labrador retriever
- Pug
- Siberian husky

2- Delete the folders associated with the remaining dog breeds in both datasets. You can also delete the folders associated with the bounding boxes in the SDD.

3- For the 7 breeds that are present in both datasets, the names might be written slightly differently (e.g. Labrador Retriever vs. Labrador). Manually rename the folders so the names match (e.g. make them both *labrador-retriever*).

4- Rename the folders to indicate that they are subsets of the original datasets (to avoid potential confusion if you later want to use them for another project). For example, *SDDsubset* and *DBIsuset*. Each of these should now contain 7 subfolders (e.g. *border-collie*, *pug*, etc.) and the names should match.



5- Zip the two folders (e.g. *SDDsubset.zip* and *DBIsubset.zip*) and upload them to your Google Drive (if you want to use Google Colab).

You can find sample code working with the SDD on the internet. If you want, you are welcome to look at these examples and use them as your starting code or use code snippets from them. You will need to modify the code as our questions are asking you to do different tasks, which are not the same as the ones in these online examples. But using and copying code snippets from these resources is fine. If you choose to use one of these online examples as your starting code, please acknowledge them in your submission. We also suggest that before starting to modify the starting code, you run them as is on your data (e.g. *DBIsubset*) to 1) make sure your dataset setup is correct and 2) to make sure you fully understand the starter code before you start modifying it.

Task I - Inspection (5%)

Look at the images in *SDDsubset* and *DBIsubset*. Do you see any systematic differences between images in *SDDsubset* and *DBIsubset*?

Task II - simple CNN (50%)

Construct a simple CNN to distinguish between the images in *SDDsubset* and the images in *DBIsubset*. For example, you can use the following architecture:

- convolutional layer (3x3, 16 filters)
- batch normalization
- convolutional layer (3x3, 16 filters)
- max pooling (2x2)
- convolutional layer (3x3, 16 filters)
- batch normalization
- convolutional layer (3x3, 16 filters)
- max pooling (2x2)
- dropout (e.g. 0.5)
- fully connected (32)
- dropout (0.5)
- softmax

If you want, you can change these specifications; but if you do so, please specify them in your submission. Use RELU as your activation function, and cross-entropy as your cost function. Train the model with the optimizer of your choice, *e.g.*, SGD, Adam, RMSProp, etc. Use random cropping, random horizontal flipping, random colour jitter, and random rotations for augmentation. Make sure to tune the parameters of your optimizer for getting the best performance on the validation set.

Plot the training, and test accuracy over the first 10 epochs. Note that the accuracy is



different from the loss function; the accuracy is defined as the percentage of images classified correctly.

Train the same CNN model again; this time, without dropout. Plot the training and test accuracy over the first 10 epochs; and compare them with the model trained with dropout. Report the impact of dropout on the training and its generalization to the test set.

Task III - ResNet Training on the DBI (15 marks):

[III.a] (10 marks) ResNet models were proposed in the “Deep Residual Learning for Image Recognition” paper. These models have had great success in image recognition on benchmark datasets. In this task, we use the ResNet-18 model for the classification of the images in the DBI dataset. To do so, use the ResNet-18 model from PyTorch, modify the input/output layers to match your dataset, and train the model for 10 epochs. Do not use the pre-trained ResNet. Plot the training and test accuracy over the first 10 epochs, and compare those with the results of your CNN model.

[III.b] (5 marks) Repeat the same task as in [III.a], but use the ResNet-34 model. Report the accuracy. Compare the accuracy obtained on the DBI dataset with the accuracy obtained on the SDD. Which is higher? Report your answer briefly, in one or two sentences.

Task IV - Fine-tuning (15 marks):

Similar to the previous task, you are supposed to use the pre-trained models (ResNet18, ResNet34, ResNet50, ResNet101, ResNet152, Swin) from torchvision or torchvision.models (within torchvision): `torchvision.models.resnet18(pretrained=True)`. Model of your choosing. Fine-tune the input/output layers and the final layer so the output matches the DBI dataset. You need to replace the final layer. **Hint:** The final layer might have a different number of classes.

This time you are supposed to use the pre-trained models and fine-tune the input/output layers on DBI training data. Report the accuracy of these fine-tuned models on DBI test dataset, and also the entire SDD dataset.

Discuss the cross-dataset performance of these trained models. Which models generalized to the new dataset better? For example, are there cases in which two different models perform equally well on the test portion of the DBI but have significant performance differences when evaluated on the SDD? Are there models for which the performance gap between the SSD and test portion of DBI are very small?

Task V - Dataset detection (15 marks):

Train a model that – instead of classifying dog breeds – can distinguish whether a given image is more likely to belong to SDD or DBI. To do so, first, you need to divide your data

into training and test data (and possibly validation if you need those for tuning the hyper-parameters of your model). You need to either reorganize the datasets (to load the images using `torchvision.datasets.ImageFolder`) or write your own data loader function. You can start from a pre-trained model (of your choice) and fine-tune it on the training portion of the dataset. Include your network model specifications in the report, and make sure to include your justifications for that choice. Report your model's accuracy on the test portion of the dataset.

Task VI - How to improve performance on SDD? (10 marks):

If our goal were to have good performance on the SDD dataset, briefly discuss how to work towards this goal in each of the following cases: (you don't need to implement these, just briefly discuss each case in 2-3 sentences)

- At training time, we have access to the training DBI dataset but not to the test portion of the SDD dataset. All we know is that we have access to the training DBI dataset. Discuss how to work towards this goal with DBI (similar to the answer you gave for Task I).
- At training time, we have access to the training DBI dataset and a small portion (e.g. 10%) of the SDD dataset. Discuss how to work towards this goal with DBI (similar to the answer you gave for Task I).
- At training time, we have access to the training DBI dataset and a small portion (e.g. 10%) of the SDD dataset. Discuss how to work towards this goal with DBI (similar to the answer you gave for Task I).

Task VII - Discussion (10 marks):

Briefly discuss how some of the solutions you proposed in Task VI can have implications in real applications. For example, consider the case where available data is from a university (e.g. a university) and the goal is to deploy trained models in a real-world setting (e.g. a university).

Summary of implementation

The train/val/test split is up to you. You can use the same split used in the sample code linked, i.e. train: 60%. validation 10%, test: 30%. But you can do other (reasonable) splits too if you want. Just specify in your report what you did.

	what we want to do	Train	Validation	Test
Task II	dog breed classification	DBI	DBI	DBI
Task III.a	dog breed classification	DBI	DBI	DBI
Task III.b	dog breed classification	DBI	DBI	SDD
Task IV	dog breed classification	DBI	DBI	both
Task V	dataset classification			